


Per-scenario multi-turn trajectories



Each panel shows one scenario. Lines correspond to Mentor (solid), Claude baseline (dashed), and GPT-5 baseline (dash-dot). Hollow markers indicate the first turn where score ≥ 1.6 .

—●— Baseline: Claude Sonnet 4.5 —●— Baseline: GPT-5 —●— Mentor (Kimi-k2-0905)